

Creating a distributed python wrapper with otwrapy

HPC and Uncertainty Treatment

PRACE Advanced Training Center, May 16-18 2018
EDF – Phimeca – Airbus Group – IMACS – CEA

Antoine Dumas, Phimeca Engineering



MAISON DE LA SIMULATION

Outline

- ④ Introduction
- ④ Basic skeleton of a wrapper
- ④ In use tools
- ④ Parallelization
- ④ Conclusion

Introduction

- ❏ A wrapper is a piece of code that creates a python interface that communicates with your code.
- ❏ A good wrapper should be able to take advantage of multi core computers and HPC clusters in order to distribute multiple evaluations of your code.
- ❏ Objective of this presentation → Show you how to create the killer distributed wrapper to efficiently carry on uncertainty studies.
- ❏ It is based on the module **otwrapy** available at **GitHub**. Initially developed at **Phimeca engineering**.
- ❏ A good working example can be found on the **otwrapy repository example**.

What makes a good wrapper ?

☐ It is distributed and avoids conflict between runs.

☐ You can use it as a script (argsparse module):

```
$ python wrapper.py -X 170 3 0.05
```

☐ It is able to run on different environments:

- ▶ Workstation
- ▶ Office made heterogeneous clusters → e.g. IPyparallel or dispy
- ▶ HPC through submission scripts → e.g. TGCC or Poincare
- ▶ Cloud solutions → e.g. Simulagora or DominoUp

☐ It catches and logs errors for easy debugging.

☐ It can either run or simply prepare runs → useful when running on clusters.

All of this might seem complex, but wrappers are repetitive and **otwrapy** is here for you !

Basic skeleton of a wrapper

- What follows assumes that you want to wrap an external code not written in Python.
- An OpenTURNS wrapper is a subclass of `ot.OpenTURNSPythonFunction` for which at least the method `_exec(X)` should be overloaded. Additionally, you can overload `_exec_sample(X)`, but with `otwrapy.Parallelizer()` there is no need to.
- If your code handles the gradient and the hessian, you can respectively overload `_gradient(X)` and `_hessian(X)`.

```
class Wrapper(ot.OpenTURNSPythonFunction):
    """Wrapper of my external code.
    """
    def __init__(self):
        """Initialize the wrapper with 4 and 1 as input and output dimension.
        """
        super(Wrapper, self).__init__(4, 1)
        # Do other stuff if necessary
    def _exec(self, X):
        """Run the model in the shell for the input vector X
        """
        pass
```

Overloading the `exec` function

☐ `_exec` is the default OpenTURNS method that executes the function on a given point.

Semantically speaking, the function is divided on three parts :

- ▶ Prepare the input parameters, e.g., create an input file.
- ▶ Run the external code, e.g., on the shell.
- ▶ Get the output parameters by parsing the output given by the code.

☐ The three steps are executed on a temporary working directory using `otwrappy.TempWorkDir`

```
def _exec(self, X):
    """Run the model in the shell for the input vector X
    """

    # Move to temp work dir. Cleanup at the end
    with otw.TempWorkDir(cleanup=True):
        # Prepare the input
        self._prepare_input(X)
        # Run the external code
        self._run_code(X)
        # Parse the output parameters
        Y = self._parse_output()

    return Y
```

Temporary working directory

- ☐ Efficiently and safely work on temporary directories with `otwrap.TempWorkDir`
 - ▶ Avoid conflict between simulations running in parallel.
 - ▶ If an exception is raised during execution, the Python interpreter should come back to the preceding current directories.
 - ▶ Cleanup the temporary working directories upon exit. Or don't if you want a full backup of the simulations.
 - ▶ Transfer necessary files needed by the external code

☐ Example:

```
import otwrap as otw
# I'm on a given dir, e.g. ~/beam-wrapper
with otw.TempWorkDir(base_temp_work_dir='/tmp', prefix='run-', cleanup=True, transfer=None):
    """
    ...
    Do stuff safely on an exclusive temporary directory and erase it afterwards
    ...
    """
    # The current working directory is something like /tmp/run-pZYpzQ

# Back on ~/beam-wrapper and /tmp/run-pZYpzQ does not exist anymore
```

Prepare the input parameters

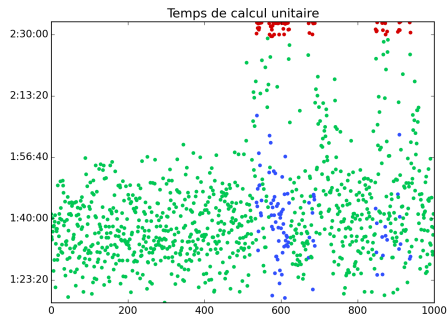
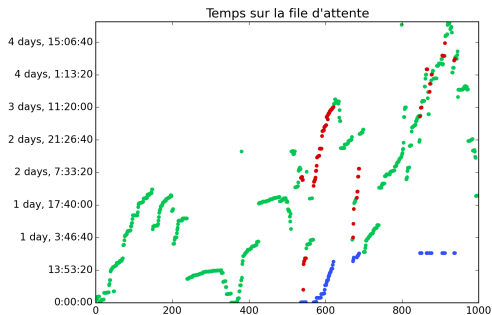
- ❏ For each simulation, your wrapper must communicate the input parameters to the external code.
- ❏ Most scientific codes use input files that describe, among other thing, the parameters of your model/simulation.
- ❏ Using OpenTURNS coupling tools, the values of the vector X are placed on an input template file that have tokens/placeholders for where the expected parameters should be.

```
def _prepare_input(self, X):  
    """Create the input file required by the code.  
    """  
    ot.coupling_tools.replace(  
        infile='input_templatefile.xml',  
        outfile='input.xml',  
        tokens=['@X1', '@X2', '@X3', '@X4'],  
        values=X)
```


Run the external code

- Most of the time this is a fairly straightforward call to an executable with an input file.
- Sometimes, it is useful to time your runtime.

```
def _run_code(self):  
    time_start = time.time()  
    ot.coupling_tools.execute('/path/to/executable -x input.xml'))  
    return time.time() - time_start
```



Parse output parameters 1/2

- ④ Common practice among scientific code is to create output files with the results of the simulation.
- ④ The output should then be parsed in order to get the output parameters of interest.
- ④ If it is a .csv file
 - ▶ `pandas.read_csv` is the fastest option, but it introduces pandas as a dependency.
 - ▶ if speed is not an issue, try `ot.coupling_tools.get_value`,
 - ▶ or `numpy.loadtxt`.

Parse output parameters 2/2

- ④ For .xml files **minidom** package from the python standard library does the trick.
- ④ If by any chance the external code returns the output parameters of interest to STDOUT, set `get_stdout=True` when calling `ot.coupling_tools.execute(...)`. (or use **`subprocess.check_output`**)
- ④ For standard binary formats, there are python interfaces to **netcdf** and **HDF5**.
- ④ Otherwise, be creative and pythonic !

```
def _parse_output(self):  
    # Retrieve output (see also )  
    xmldoc = minidom.parse('outputs.xml')  
    itemlist = xmldoc.getElementsByTagName('outputs')  
    Y = float(itemlist[0].attributes['Y1'].value)  
  
    return [Y]
```

Managing data backups

- ☐ Uncertainty studies tend to be expensive in computational time, it is then in your best interest to backup your simulation results !
- ☐ otwrapy has two useful functions to do so: `otwrapy.dump_array` and `otwrapy.load_array`
- ☐ They are faster than a simple `pickle.dump` and `pickle.load` (because they use `protocol=2`)
- ☐ They offer the possibility to compress data with the `gzip` library. If the extension is `'pklz'`, it compresses by default.
- ☐ Advice: Convert your `ot.Sample` to a `np.array` before dumping. An `np.array` is lighter !
- ☐ Example: dump and compress

```
import otwrappy as otw
otw.dump_array(np.array(X), 'InputSample.pklz', compress=True)
```

- ☐ ... and load

```
import otwrappy as otw
import openTurns as ot
X = otw.load_array('InputSample.pklz')
X = ot.Sample(X)
```

Catch exceptions when your code fails

- ☐ In order to catch exceptions use the decorator `otwrapy.Debug()` !
- ☐ It encloses what happens inside a function into a try/catch structure and logs Exceptions when they are raised.
- ☐ Useful when your wrapper is not used on an interactive environment like IPython or a Jupyter notebook.

```
import otwrappy as otw
class Wrapper(ot.OpenTURNPythonFunction):
    @otw.Debug('wrapper.log')
    def _exec(self, X):
        #Do stuff
        return Y
```

Creating a CLI for your wrapper

- ☐ A command line interface allows you to run your wrapper in detached mode, e.g., through submission scripts on HPC clusters.
- ☐ The **argparse** library might seem complicated, but they have a great cookbook and there are good chances that a simple copy/paste will be enough.
- ☐ Take a look at the **beam wrapper** for an example of a CLI interface

```
if __name__ == '__main__':  
    import argparse  
    parser = argparse.ArgumentParser(description="Python wrapper example.")  
    parser.add_argument('-X', nargs=3, metavar=('X1', 'X2', 'X3'),  
                        help='Vector on which the model will be evaluated')  
    args = parser.parse_args()  
  
    model = Wrapper(3, 1)  
    X = ot.NumericalPoint([float(x) for x in args.X])  
    Y = model(X)  
    dump_array(X, 'InputSample.pkl')  
    dump_array(Y, 'OutputSample.pkl')
```

- ☐ You can then execute your code from the command line :

```
python wrapper.py -X 170 3 0.05
```



Parallelizing the wrapper

- ☐ Uncertainty studies fall into what we call *embarrassingly parallel* (or *pleasantly parallel*) patterns → Repeat similar non communicating tasks over and over.
- ☐ Good news, this means that they are very simple to parallelize.
- ☐ But don't bother... just let the magic happen with `otwrapy.Parallelizer()` !!

```
import otwrapy as otw
from otwrapy.examples.beam import Wrapper
parallelized_beam_wrapper = otw.Parallelizer(Wrapper())
```

Distributing calls on clusters or the cloud

- ❏ But what if you want to distribute your wrapper calls on the cloud or on a cluster ?
- ❏ `otw.Parallelizer` is no longer the way to go, for the moment. . .
- ❏ You can manage to make an heterogeneous office cluster with `IPyparallel` or `dispy`
- ❏ For clusters and the cloud, rely on a good CLI interface of your wrapper and distribute your calls through submission scripts or cloud APIs (e.g., `Simulagora` or `Domino`)

Conclusion

- ☐ Take away message : Making a wrapper is all about preparing the input, executing the code and parsing the output on isolated working directories. Don't forget, in a multi-core era you don't have a choice, make your wrapper distributed !
- ☐ By creating a CLI of your wrapper, you can easily distribute your calls on a cluster or on cloud platforms.
- ☐ It is important to protect your wrapper with `otw.Degbug()` so that you can have a traceback of raised Exceptions.
- ☐ `ot.PythonFunction()` is a simpler alternative to `ot.OpenTURNPythonFunction()`, but you loose the ability to parameterize your wrapper when instantiating it.
- ☐ `otwrapy` is here for you ! Use it to avoid code boilerplate or as a simple cookbook.

Thank you for your attention



Antoine Dumas

dumas@phimeca.com

Github : [otwrapy](#)