

# Low-rank in Uncertainty Management

## Efficient linear algebra and an overview of the $\mathcal{H}$ -matrix framework

Kieran Delamotte

IMACS

May, 2020



First encounter with low rank

Large and dense systems

$\mathcal{H}$ -matrices

Applications

Conclusions & perspectives

## First encounter with low rank

Large and dense systems

$\mathcal{H}$ -matrices

Applications

Conclusions & perspectives

## Key ideas for fast computation

Uncertainty quantification (UQ) typically requires a high number of simulations using basic linear algebra operations. Then two things come in handy :

- fast linear algebra decomposition ( $QR$ ,  $LL^T$ , ...)
- efficient data structures

If many simulations are 'similar' there is most likely a low-rank property somewhere to be exploited. Low-rank property in linear algebra leads to fast computations AND efficient structures !

We shall present two efficient methods :

- random linear algebra for medium-sized problems (on a simple example !)
- the  $\mathcal{H}$ -matrix framework (much more detailed)

## Model problem

Suppose we are interested in eigenvalues (and/or eigenvectors) of a covariance matrix of the form  $C_X = X^T X$  where  $X$  is of size  $m \times n$ . Several methods :

- eigenvalue decomposition of  $C_X$  :  $\mathcal{O}(n^3)$  operations and conditionning of normal equations is bad !
- SVD decomposition of  $X$  : (better, still expensive)  $\mathcal{O}(mn^2)$  operations
- what else ?

## Randomness and linear algebra

**Basic idea :** using random test matrix  $\Omega$  and the fact that any orthogonal basis of  $\text{ran}(Y)$  with  $Y = A\Omega$  is a good approximate basis for  $\text{ran}(A)$ .  $Y$  has a less columns than  $A$  so  $QR$  is cheaper !

---

**Algorithm 1** A simple random QR decomposition

---

**Require:** A matrix  $\mathbf{A}$  of size  $m \times n$ , an oversampling parameter  $l$ .

**Ensure:** An orthogonal basis  $\mathbf{Q}_Y$  of  $\text{ran}(\mathbf{A})$ .

- 1: Draw a random matrix  $\mathbf{W}$  of size  $n \times l$ .
  - 2: Form product  $\mathbf{Y} : \mathbf{Y} = \mathbf{AW}$ .
  - 3: Form the  $QR$  decomposition of the matrix  $\mathbf{Y} : \mathbf{Y} = \mathbf{Q}_Y \mathbf{R}$ .
  - 4: **return**
-

## A random SVD

---

### Algorithm 2 A simple random SVD decomposition

---

**Require:** A matrix  $\mathbf{A}$  of size  $m \times n$ , an orthogonal matrix  $\mathbf{Q}$  tq  
 $\mathbf{A} \simeq \mathbf{Q}\mathbf{R}$ .

**Ensure:** An approximate SVD decomposition of  $\mathbf{A}$ ,  $\mathbf{A} \simeq \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$ .

Construct the projection matrix  $\mathbf{B} = \mathbf{Q}^T \mathbf{A}$ .

2: Form the SVD of  $\mathbf{B}$  :  $\mathbf{B} = \tilde{\mathbf{U}}\mathbf{\Sigma}\mathbf{V}^T$ .

Build the matrix  $\mathbf{U} = \mathbf{Q}\tilde{\mathbf{U}}$ .

4: **return**

---

## A numerical example

Say  $X \in \mathbb{R}^{m \times n}$  is a discrete approximate of a random process  $X(t, \omega)$  over  $[-1, 1]$  with exponential covariance function  $C(s, t) = e^{-|t-s|}$ . We are still interested in eigenvalues of  $C_X$ . In that case, we can prove that  $\lambda_n/\lambda_1 = \mathcal{O}(n^{-2})$ . In the following example we choose  $m = 6400$  and  $n = 1000$ .



# Eigenvalues estimates




FIGURE – Eigenvalues decay.




FIGURE – Eigenvalues decay - the first 10.

## Case summary

test matrix # cols.	elapsed time (s)	$\sum \lambda_n$	rel. error
5	0.021	4950345.8	0.253
10	0.027	5894250.4	0.111
20	0.038	6309314.4	0.048
40	0.060	6478969.3	0.023
80	0.068	6556588.8	0.011
160	0.131	6598333.6	0.005

TABLE – Example summary for random SVD

- For reference :  $\text{trace}(C_X) = 6632269.0$
- elapsed time for full SVD : 0.663s
- elapsed time for *numpy.linalg.eigs* : 107.307.

## Partial conclusion

- many articles in the literature,
- random SVD is implemented in OPENTURNS,
- generally behave better than classical methods,
- what if the probleme is larger ?

First encounter with low rank

Large and dense systems

$\mathcal{H}$ -matrices

Applications

Conclusions & perspectives

# Introduction

Several applications in statistics lead to large and dense matrices :

- Kriging ;
- applications using large covariance matrices (ex : random Gaussian sampling).

## Kriging

- Symetric Positive Definite (SPD) matrix ;
- Many solves ;
- Usually,
  - Cholesky decomposition  $M = LL^T$  ;
  - Forward/Backward substitutions.

# Kriging

## in a nutshell

Let  $Z : \mathbb{R}^d \mapsto \mathbb{R}$  be some random process, stationary of order 2. We have  $N$  observation points  $\{x_i \in \mathbb{R}^d / i = 1, \dots, N\}$  with values  $Z(x_i)$ . Let assume that the covariance of these points is known and given by a matrix  $K \in \mathbb{R}^{N \times N}$  with

$$K_{ij} = \text{Cov}(Z(x_i), Z(x_j))$$

An estimation of the mean trajectory is a linear interpolation  $\tilde{Z}(x_0)$  of  $Z$  at  $x_0 \in \mathbb{R}^d$  :

$$\tilde{Z}(x_0) = \sum_{i=1}^N \alpha_i(x_0) Z(x_i)$$

The weights  $\alpha_i$  are solution of the linear system

$$K\alpha = K_0 \quad \text{where } (K_0)_i := \text{Cov}(Z(x_0), Z(x_i))$$

# Kriging

## writing the systems

- $K$  is SPD (covariance matrix);
- usually  $K$  is not known exactly :
  - modelled as a convolution matrix of a kernel  $k : \mathbb{R}^d \mapsto \mathbb{R}$  :

$$K_{ij} := k(x_i, x_j)$$

- what is  $k$  ?
  - Exponential :

$$k(x, y) = e^{-|x-y|/\lambda}$$

- Gaussian :

$$k(x, y) = e^{-|x-y|^2/(2\lambda^2)}$$

- Quadratic :

$$k(x, y) = \left(1 + \frac{|x-y|}{2\lambda}\right)^{-2}$$

- $N$  can be large. For instance, every node in a FEM discretization ;
- The interpolation is sought at many points  $x_0$ .

# Kriging

## solving the system

- $K$  is ill-conditioned, many RHS : direct solver ;
- $K$  is SPD : Cholesky.

## Complexity (LAPACK)

- Cholesky factorization (DPOTRF) :  $1/3N^3 + 1/2N^2 + 1/6N$
- Solving (DPOTRS) :  $N_{\text{RHS}} \times 2N^2$
- Storage :  $8N^2/2$

Need of a fast direct solver :  $\mathcal{H}$ -matrix framework.



## Toy model : exponential kernel in 1D

Let  $X_N$  be a uniform discretization of  $[0, 1]$  with the discretization step  $h = \frac{1}{N-1}$  :

$$X_N = \{0 = x_1, \dots, x_N = 1\}$$

The correlation length  $\lambda$  is set to  $\lambda := 5h$  and the exponential kernel in 1D reads as

$$K_\lambda(x_i, x_j) = e^{-|x_i - x_j|/\lambda}$$

# Toy model : exponential kernel in 1D

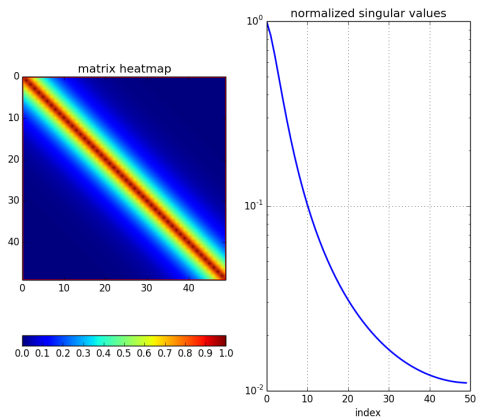


FIGURE – the covariance matrix  $K_\lambda([0, 1], [0, 1])$

# Toy model : exponential kernel in 1D




FIGURE – small extra-diagonal block :  $K_\lambda([0, 0.2], [0.8, 1])$

# Toy model : exponential kernel in 1D




FIGURE – large extra-diagonal block :  $K_\lambda([0, 0.5], [0.5, 1])$

# Toy model : exponential kernel in 1D

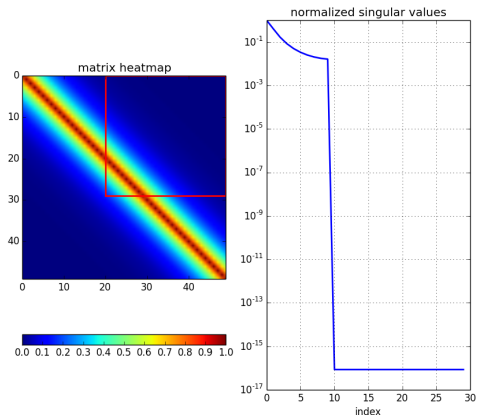


FIGURE – taking a part of the diagonal :  $K_\lambda([0, 0.6], [0.4, 1])$

## Toy model : exponential kernel in 1D

Whenever  $x_i$  and  $x_j$  are in disjoint sets the kernel reads as a separated one. For instance ; if  $x_i > x_j$  then

$$K_\lambda(x_i, x_j) = e^{(x_j - x_i)/\lambda} = e^{x_j/\sqrt{\lambda}} e^{-x_i/\sqrt{\lambda}},$$

which is of rank 1.

First encounter with low rank


Large and dense systems

$\mathcal{H}$ -matrices

Applications

Conclusions & perspectives

# Admissibility gives low-rank



Usual condition :

$$\min(\text{diam}(B_t), \text{diam}(B_s)) \leq \eta \text{dist}(B_t, B_s) \quad (\text{admissibility condition})$$

The separation condition  $R = 0$  gives the **HODLR**(**H**ierarchically **O**ff-**D**agonal **L**ow-**R**ank) structure described by the 1D toy model.



## Low-rank approximation : compression techniques

- SVD :  $M \approx U\Sigma V^H$ 
  - Rank and precision controlled ;
  - **Costly**  $\mathcal{O}(4m^2n + 8mn^2 + 9n^3)$  (hyp :  $m > n$ )
- Existence of cross approximations : row/col. extraction (Goreinov & Tyrtysnikov '97) ;
- Gaussian/LU rank-revealing scheme known as **Full Cross Approximations** :

---

```
1: while  $\|M\| \geq \varepsilon \|M_0\|$  : do
2:    $\text{rank}(M) \leftarrow \text{rank}(M) + 1$ 
3:   Find the coefficient  $M_{i^*j^*}$  so that  $M_{i^*j^*} = \max_{i,j} |M_{ij}|$ ,  $\alpha =$   
      $M_{i^*j^*}$ 
4:    $M \leftarrow M - \frac{1}{\alpha} M(:, j^*) M(i^*, :)$ 
5: end while
```

---

## Variants

The fast determination of the pivot is the main idea of all fast algorithms. Key points to speed up the full cross approximation :

- **Partially pivoted Cross Approximation.**
  - We seek the largest pivot over a column and/or a row in  $\mathcal{O}(m)$  instead of  $\mathcal{O}(m^2)$  operations.
  - Only the modified coefficients of the remainder are computed at each step.
- **Adaptive CA algorithm** : a fast (linear) estimation of the remainder.
- **ACA+** and other variants use other heuristics.
- Trade-off between robustness (SVD) and efficiency (ACA/ACA+) : computations from  $\mathcal{O}(m^2n + mn^2)$ (SVD) to  $\mathcal{O}(mnr)$ (fullCA) to  $\mathcal{O}((m+n)r^2)$ (ACA).

## Space partitioning : clustering

- Use of **bounding boxes** : easier to handle than point clouds ;
- Recursive splitting strategy (**Divide and Conquer strategy**) thanks to **nested bisection** :
  - geometric : the box is split in two halves along the largest axis ;
  - median : each half contains roughly the same number of unknowns ;
  - others (PCA,...).
- Split the boxes until each box contains a fixed small number of unknowns ;
- Result : **binary tree**.

## Space partitioning : clustering




FIGURE – Illustration of the geometric clustering on a cone-sphere.

## Space partitioning : clustering




FIGURE – Illustration of the geometric clustering on a cone-sphere.

## Space partitioning : clustering




FIGURE – Illustration of the geometric clustering on a cone-sphere.

## Space partitioning : clustering




FIGURE – Illustration of the geometric clustering on a cone-sphere.

## Blockclustering : Clustering & Admissibility

It is a quad-tree whose nodes are matrix blocks and the leaves are admissible (or small) blocks; a block is split in a  $2 \times 2$  block structure.





TABLE – Blockclustering and the geometry



# The $\mathcal{H}$ -matrix structure

A  $\mathcal{H}$ -matrix is a **quadtree** (with Binary Space Partitioning) :

- Internal nodes : subdivided  $\mathcal{H}$ -matrix ;
- Leaves :
  - admissible block : large & low-rank ;
  - inadmissible block : dense & full rank, but small.

## Remarks

- Only the leaves carry data ;
- Big admissible blocks ( $10^4 \times 10^4$  and more)
- Small (and few) inadmissible blocks ( $100 \times 100$ ).

Each admissible block is compressed with a fast method thus determining a numerical rank with a prescribed relative error  $\varepsilon$ .

## Operations : three kinds

- $\mathcal{H}$  -**BLAS1&2** : Assembly, AXPY, GEMV  
Simple. Operating only on leaves.
- $\mathcal{H}$  -**BLAS3** : GEMM, TRSV  
More involved. Operations at many different levels of the same sub-tree.
- $\mathcal{H}$  -**LAPACK** : Inverse,  $LU$ ,  $LL^T$ .  
Uses BLAS2 and BLAS3 operations, harder to implement in parallel.

### Base operations

All operations use BLAS/LAPACK. Typical subroutines are : SVD, QR, LU, TRSV, GEMM and GEMV.


# Addition

- *Usually* structures must be the same ;
- 3 different base cases :
  - sum of two dense matrices (usual dense operation) ;
  - sum of a dense and a low-rank matrix ;
  - **sum of two low-rank matrices.**

## Useful pointers

- Adding two low-rank matrices : unknown final rank
- Costly operation !

# Multiplication




## Main issues

Let  $C = A \times B$  the product of 2  $\mathcal{H}$ -matrices.

- What is the 'best' structure of  $C$ ?
- Uniqueness?
- What if it is imposed?

# Multiplication




## Main issues

Let  $C = A \times B$  the product of 2  $\mathcal{H}$ -matrices.

- What is the 'best' structure of  $C$ ?
- Uniqueness? (here the literature definition)
- What if it is imposed?

# Multiplication




## Main issues



Let  $C = A \times B$  the product of 2  $\mathcal{H}$ -matrices.

- What is the 'best' structure of  $C$ ?
- Uniqueness?
- What if it is imposed? (e.g. in a  $LL^T$  factorization)

## Multiplication : exemple 1




# Multiplication : exemple 1






## Multiplication : exemple 2



## Multiplication : exemple 2



# Cholesky factorization

Recall : recursive block splitting based on the  $2 \times 2$  block structure.

$$\begin{bmatrix} A_{11} & A_{21}^T \\ A_{21} & A_{22} \end{bmatrix} = \begin{bmatrix} L_{11} & 0 \\ L_{21} & L_{22} \end{bmatrix} \times \begin{bmatrix} L_{11}^T & L_{21}^T \\ 0 & L_{22}^T \end{bmatrix}$$

$$A_{11} = L_{11} L_{11}^T$$

$$A_{21} = L_{21} L_{11}^T$$

$$A_{22} = L_{21} L_{21}^T + L_{22} L_{22}^T$$

## Cholesky factorization

Recall : recursive block splitting based on the  $2 \times 2$  block structure.

$$\begin{bmatrix} A_{11} & A_{21}^T \\ A_{21} & A_{22} \end{bmatrix} = \begin{bmatrix} L_{11} & 0 \\ L_{21} & L_{22} \end{bmatrix} \times \begin{bmatrix} L_{11}^T & L_{21}^T \\ 0 & L_{22}^T \end{bmatrix}$$

$$A_{11} = L_{11} L_{11}^T \quad (\text{recursive call})$$

$$A_{21} = L_{21} L_{11}^T$$

$$A_{22} = L_{21} L_{21}^T + L_{22} L_{22}^T$$

# Cholesky factorization

Recall : recursive block splitting based on the  $2 \times 2$  block structure.

$$\begin{bmatrix} A_{11} & A_{21}^T \\ A_{21} & A_{22} \end{bmatrix} = \begin{bmatrix} L_{11} & 0 \\ L_{21} & L_{22} \end{bmatrix} \times \begin{bmatrix} L_{11}^T & L_{21}^T \\ 0 & L_{22}^T \end{bmatrix}$$

$$\begin{aligned} A_{11} &= L_{11} L_{11}^T \\ A_{21} &= L_{21} L_{11}^T \quad (\text{upper triangular solve}) \\ A_{22} &= L_{21} L_{21}^T + L_{22} L_{22}^T \end{aligned}$$

# Cholesky factorization

Recall : recursive block splitting based on the  $2 \times 2$  block structure.

$$\begin{bmatrix} A_{11} & A_{21}^T \\ A_{21} & A_{22} \end{bmatrix} = \begin{bmatrix} L_{11} & 0 \\ L_{21} & L_{22} \end{bmatrix} \times \begin{bmatrix} L_{11}^T & L_{21}^T \\ 0 & L_{22}^T \end{bmatrix}$$

$$A_{11} = L_{11}L_{11}^T$$

$$A_{21} = L_{21}L_{11}^T$$

$$A_{22} = L_{21}L_{21}^T + L_{22}L_{22}^T \quad (\mathcal{H}\text{-matrix GEMM then recursive call})$$

# Computation details

## Typical issues

The  $\mathcal{H}$ -matrix algorithms are not nice for the hardware :

- Very small operations ;
- Oddly-shaped matrices : "Tall & skinny" ;
- High memory band.

## Observations

- Most (70-80%) of the time spent in :
  - QR decompositions of T&S matrices ;
  - SVD decompositions of small matrices.
- BLAS implementations cannot reach the peak performance ;
- Very high memory bandwidth.

## Complexity estimates

### Useful pointers

- Most complexity estimates assume a fixed upper bound  $k$  for the low-rank matrices involved ;
- The structure of the matrix (as represented by a tree) is important as well : a large depth with small blocks is typically a bad omen.

### Common operations

For a matrix size of  $N \times N$  with the previous assumptions :

- assembly : time and storage is in  $\mathcal{O}(kN \log N)$  ;
- addition :  $\mathcal{O}(k^2 N \log N)$  operations ;
- multiplication and Cholesky factorisation :  $\mathcal{O}(k^3 N \log^3 N)$  operations ;
- in practice a  $\mathcal{O}(N \log^2 N)$  complexity is observed.



First encounter with low rank

Large and dense systems

$\mathcal{H}$ -matrices

**Applications**

Conclusions & perspectives

## Recall the 1D toy model ?

- exponential kernel :  $K(s, t) = e^{-|s-t|/\lambda}$
- prescribed relative error  $\varepsilon = 10^{-4}$  ;
- **HODLR** admissibility ;
- 1D exponential kernel : provably rank-one when **HODLR**.
- 'exact model' is the second Hermite function  
 $\psi_2(x) = (2x^2 - 1)e^{-\frac{1}{2}x}$  ;
- input data as a gaussian distribution ;
- here : one iteration of optimization loop, correlation length  $\lambda = 0.01$ .

## Recall the 1D toy model ?




FIGURE – Lower part of the covariance matrix : rank map.

- matrix size  $1000 \times 1000$  ;
- compression ratio :  $\approx 12\%$  (small case !)

## Recall the 1D toy model ?




FIGURE – Cholesky factor : rank map.

- matrix size  $1000 \times 1000$  ;
- compression ratio :  $\approx 12\%$  (small case !)

## Recall the 1D toy model ?





FIGURE – Surface response using kriging


- compression ratio :  $\approx 12\%$  (small case!);
- same response with LAPACK and HMAT solver;
- maximum absolute error between two approximates :  $1.55 \times 10^{-15}$ .

# Performances - computational time

## the exponential kernel in 1D



(a) Matrix assembly




(b) Matrix factorization


FIGURE – Computational time vs matrix size

# Performances - memory

## the exponential kernel in 1D



(a) Matrix assembly



(b) Matrix factorization

FIGURE – Memory vs matrix size

## 3D exemple : a cantilever beam

- The vertical deflection  $y$  of a cantilever beam's free end of fixed length  $L$  reads as

$$y = \frac{FL^3}{3EI},$$

where :

- $E$  is the Young modulus ;
- $F$  is the load ;
- $I$  is the moment of inertia.
- Input* variables  $x = (F, E, I)$  are assumed random ;
- Variable of interest (*output*) is the deflection  $y$  estimated thanks to the model  $\mathcal{M}$

$$\mathcal{M} : x \mapsto y$$

- Building a metamodel  $\tilde{\mathcal{M}}$  through Kriging and optimization loop for parameters : one  $\mathcal{H}$ -matrix at each iteration to treat the covariance matrix associated with a specified covariance model.



The chosen covariance model  $K_\lambda(s, t)$  is the following tensor product :

$$e^{-(s_1-t_1)/\lambda_1} e^{-(s_2-t_2)/\lambda_2} e^{-(s_3-t_3)/\lambda_3}.$$

- Let  $\lambda_1 = 3.96528$ ,  $\lambda_2 = 5.8237$  and  $\lambda_3 = 9.0679$  be the starting coefficients for the optimization.
- Degrees of freedom to be clustered within the  $\mathcal{H}$ -matrix framework :

"Young modulus"; "Load"; "Inertia"

3.6258375026+07; 4.797336026+04; 3.5171332517+02

3.1382130227+07; 3.350317520+04; 3.5324327901+02

... ..

# Results




FIGURE – Input data :  $5 \times 10^4$  entries.

- covariance matrix of size  $5.10^4 \times 5.10^4$ , symmetric and double precision : full size in memory is 10Go.

## Results




FIGURE – Lower part of the covariance matrix : rank map.

- memory : 167Mo (compression ratio : 1.67%);
- assembly time : 11.83s

## Results




FIGURE – Cholesky factor : rank map.

- memory : 263Mo (compression ratio : 2.63%);
- Cholesky time : 17.84s

First encounter with low rank

Large and dense systems

$\mathcal{H}$ -matrices

Applications

Conclusions & perspectives

## Summing up the $\mathcal{H}$ -matrices method

- Three key components for assembling a  $\mathcal{H}$ -matrix :
  - The clustering of degrees of freedom (e.g. geometric) ;
  - An admissibility condition = which block to compress ;
  - A fast on-the-fly algorithm to assembly low-rank admissible blocks.
- An algebra on  $\mathcal{H}$ -matrices :
  - Multiplications and additions of  $\mathcal{H}$ -matrices ;
  - Fast factorization of an  $\mathcal{H}$ -matrix with the same structure :
    - Fast direct solver ;
    - Good preconditioner for iterative solver.

# Conclusion

- The  $\mathcal{H}$ -matrix framework is an enabler for many statistics problems ;
- Sequential solver is freely available through OpenTurns ;
- Efficient parallel solver through licensing (contact Airbus & Imacs).